

# TCP Congestion Control -- Overview

Kameswari Chebrolu

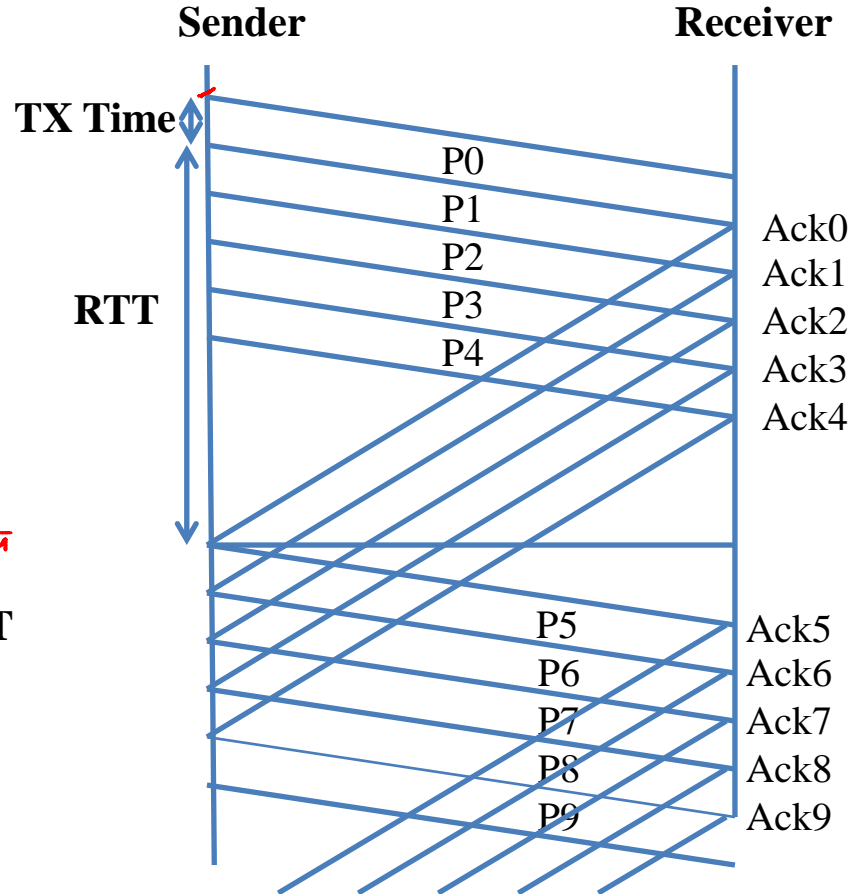
Seminal Paper: Congestion Avoidance and Control  
by Van Jacobson and Michael J. Karels

# Recap: TCP Services

- Multiplexing/Demultiplexing
- Reliable point-to-point data transfer
- Full-duplex
- Congestion control
- Flow control

Sliding  
window  
protocol

# Recap: Sliding Window

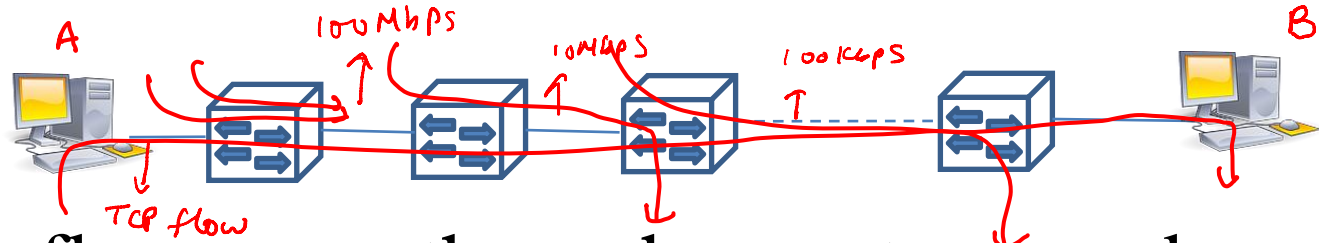


Throughput  $\sim (W * MSS) / RTT$

*Handwritten annotations:*  
- A red circle around  $W$  with an arrow pointing to it from the text "wind. size pkts".  
- A red arrow pointing up to  $MSS$ .

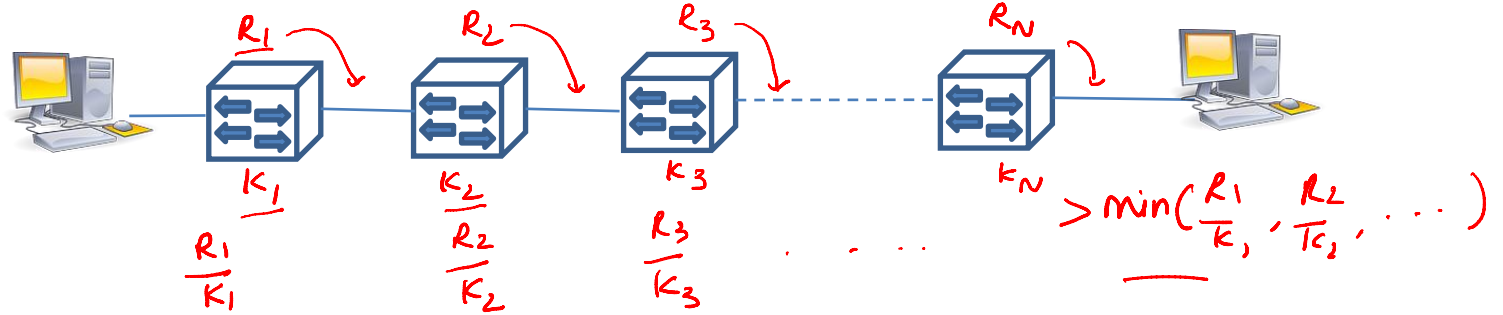
ignore tx delays if W large

# Congestion Control: Problem Statement



- Many flows pass through a router; number varies with time
- Flows can be TCP or UDP
- The link capacities of the routers are different
- End Result: Throughput achieved by a given flow function of many factors

# Congestion Control: Challenge



- Need to estimate  $W$  (of sliding window) such that each flow gets its fair share
  - Estimate small  $\rightarrow$  underutilization; Estimate large  $\rightarrow$  Congestion
- $W$  will vary over time
- Congestion Control: Preventing sources from sending too much data too fast and thereby ‘congest’ the network


# Sliding Window Protocol

- Roughly, idea translates to the following:
- View network as a pipe
- Determine the capacity of the pipe (Bandwidth-delay product)
- Fill the pipe with data
- As you remove one packet from the pipe, add another
  - ACKs help clock out data (Self Clocking)

# 3 Steps

- Getting to Equilibrium<sup>↗</sup>
- Conservation at equilibrium<sup>→</sup>
  - Don't put new packet unless old one is removed
- Adapting to Path Dynamics

# Summary

- Congestion Control is a complex problem
- Need to implement it in the context of the sliding window protocol
  - Self clocking  is a useful feature
  - Need to determine and adapt  $W$  (window size) such that you don't underutilize bandwidth or congest the network
- Ahead: Actual details 